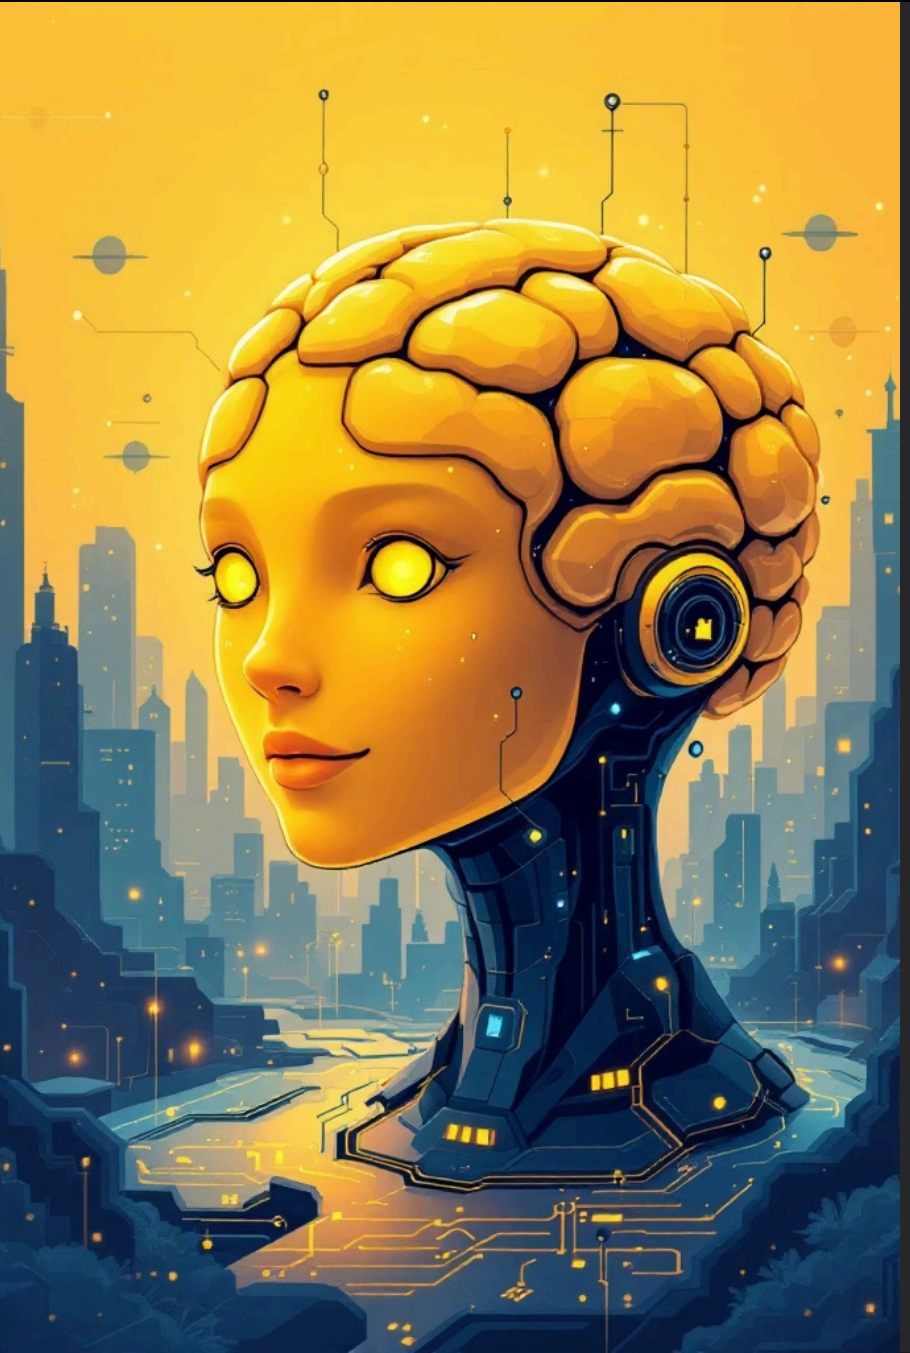

Local-First AI for Sensitive Legal Data

This 14-week simulation challenges you as a Lead AI Architect to automate departments weekly, focusing on practical, cost-effective, and secure AI solutions.

Made with GAMMA

The Program: A 14-Week AI Simulation

Real-World Automation

This isn't theory; it's hands-on. We build systems that deliver immediate, tangible savings in time and money.

Zero-Barrier Tech

We prioritize free-tier tools and implement "Zero-Cost Hacks" to ensure accessibility for all participants.

Build in Public

Every Sunday, you'll ship your code and share a demo video, fostering accountability and community learning.

Made with GAMMA

WEEK 02: The Local "Privacy" LLM

Mission Critical: Sensitive Legal Contracts

Your client has a strict requirement: sensitive legal contracts must never touch the cloud. Your mission is to deploy an offline AI brain that runs entirely on a laptop, ensuring maximum data privacy and security.

Made with GAMMA

Legal & Data Security Automation: Offline Knowledge Chat

The Problem: Data Sovereignty

"Legal documents cannot leave the company laptop." This constraint drives the need for an air-gapped solution.

The Build: Private AI Brain

You will build a Local LLM capable of PDF ingestion and private Q&A, creating a secure chat interface.

The Lesson: Privacy-First AI

This week teaches the deployment of local models and the principles of privacy-first AI architecture.

Made with GAMMA

The Tech Stack: Empowering Local AI

Ollama

A powerful tool for running large language models locally, bringing advanced AI capabilities to your desktop.

Docker

Enables containerization, ensuring consistent and isolated environments for your AI applications.

AnythingLLM (Desktop)

Provides a user-friendly interface for interacting with your local LLMs, making private AI accessible.

Made with GAMMA

The Zero-Cost Hack: Overcoming Hardware Limitations

Leveraging Google Colab for Remote-Local Access

If your laptop's performance is a bottleneck, we implement the "Remote-Local" hack. Google Colab hosts the Ollama tunnel via Ngrok, allowing you to chat with your private LLM remotely and for free.

This setup ensures that data never leaves your secure control, even when leveraging remote computational resources.

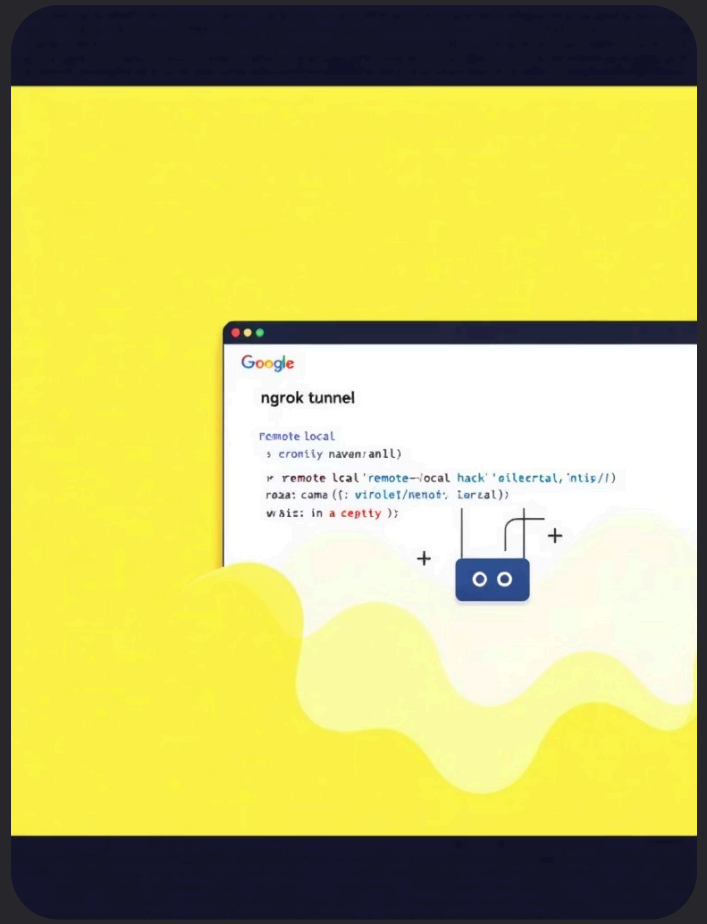

Made with GAMMA

Deliverable: A Private, Secure Chat Interface

By the end of Week 2, you will have successfully deployed a fully functional, private chat interface. This system is designed to handle sensitive legal data, operating entirely offline or via a securely tunneled connection, without ever exposing information to external servers.

Secure

Ensuring compliance with stringent privacy requirements.

Chat-Enabled

An intuitive interface for interacting with legal documents.

Offline Ready

Designed for environments without internet access.

Made with GAMMA

The Architectural Implication: Air-Gapped Power

As an AI Architect, your challenge is to marry the power of advanced LLMs with the ironclad security of an air-gapped system. This involves careful consideration of deployment, data handling, and user interaction to maintain absolute privacy.

Made with GAMMA

Reel Angle: "This AI Never Touches the Internet."

This week's solution highlights a critical capability: providing advanced AI functionality where data privacy is paramount. The "reel angle" emphasizes the core value proposition – a powerful AI that operates with complete data isolation, perfect for the most sensitive legal and enterprise data.

Made with GAMMA